

Referee Report on “Evolving Reputation for Commitment: Understanding Inflation and Inflation Expectations”

This paper develops a theory of inflation and inflation expectations in which private agents learn over time about the central bank’s commitment to its inflation target. The key state variable is reputation, defined as agents’ belief that the policymaker can credibly commit to future policy. Inflation outcomes that deviate from announced targets update this belief, which in turn feeds back into expectations and policy effectiveness.

The authors embed this learning mechanism into a standard New Keynesian model with two types of policymakers, one committed and one opportunistic, whose actions are imperfectly observed. They show that optimal policy internalizes how current inflation choices affect future beliefs, leading the committed policymaker to strategically manage reputation. They solve the resulting dynamic game using recursive principal–agent formulation.

Quantitatively, the model is calibrated to U.S. data using Survey of Professional Forecasters (SPF) inflation expectations to extract latent reputation and shock processes. The estimated reputation falls sharply during the 1970s and recovers after the Volcker disinflation. The model is able to match observed inflation dynamics and predicts a time-varying sensitivity of long-run inflation expectations to inflation surprises, which the authors validate using SPF forecast revisions.

- A central claim of the paper is that high reputation compresses the equilibrium policy gap between committed and opportunistic policymakers, while low reputation widens it, because the committed policymaker has stronger incentives to accelerate private-sector learning when reputation is low. I think the paper would benefit from a clearer discussion of which assumptions and which forces drive this monotone relationship. While the paper provides intuition and numerical illustrations supporting the pattern, the direction of the comparative statics is not obvious a priori, because there are plausible countervailing effects.

In particular, the equilibrium mapping from reputation to the policy gap is poten-

tially theoretically ambiguous. When reputation is very low, beliefs may respond only weakly to typical inflation realizations, so inducing meaningful updating may require large policy separation, which could be costly for the committed type. Moreover, if the opportunistic type were forward-looking, it would have incentives to mimic the committed type to preserve or rebuild reputation, which would tend to compress the policy gap precisely in low-reputation states. Conversely, when reputation is very high, deviations may have relatively small reputational consequences at the margin, which could in principle allow a forward-looking opportunistic type to move closer to its preferred policy without immediately triggering large belief revisions—potentially widening the policy gap. More generally, one could imagine non-monotonic patterns, with the largest gaps arising at intermediate reputation levels where incentives to manipulate beliefs are strongest.

To be clear, I am not asking the authors to solve the full model with a forward-looking opportunistic policymaker. Rather, given how prominently the monotone logic is used in interpreting the Volcker episode and the time-varying forecast revision evidence, I think the paper should more explicitly explain why this monotone result emerges in their environment and which elements (e.g., opportunistic myopia, expectation feedback strength, calibration) are doing the work.

- The paper develops nice machinery to solve a tricky dynamic learning problem. A key challenge, however, is to convince the reader that this machinery is necessary for understanding observed inflation episodes. At present, the quantitative exercise falls somewhat short of doing so. In particular, the paper extracts latent states and shocks using the full strategic model and then evaluates a “naive policymaker” counterfactual by feeding those same filtered states into an alternative policy problem. While this is a useful policy counterfactual, it does not establish that the naive model fails empirically on its own terms. Because the filtering step embeds the strategic learning structure, the naive model is never asked to explain the data using its own implied belief dynamics and state evolution. A more convincing exercise would be to re-run the filtering exercise under the naive model and show that it cannot jointly account for SPF expectations and inflation dynamics.
- It would be useful to clarify whether the data actually calls for the specific commitment vs. opportunistic structure, as opposed to a more reduced-form model of unknown policymaker preferences. For example, consider an alternative in which policymakers are always “opportunistic” (i.e., reoptimize each period) but come in two unobserved types that differ in a preference/target parameter. Private agents would learn about this hidden parameter from inflation realizations, which could also generate time-varying

long-run expectations and time-varying forecast-revision sensitivities.

- The paper argues that a key measure of success is its ability to match observed inflation without directly targeting inflation, relying instead on survey expectations. Given the emphasis placed on this result, I think a more detailed discussion of the model's calibration would be useful. In particular, several parameters that are important to inflation dynamics, such as the output weight, the output target, and the parameters governing initial and inherited reputation, are set exogenously and are not tightly disciplined by data. How sensitive is the model's ability to match inflation to these choices? If the fit depends importantly on them, it would be helpful to discuss whether there are alternative ways of pinning them down more formally (for example, via indirect inference or by targeting additional moments), rather than treating them as assumptions.